

Solution scope

We propose to use Jupyter Notebook with proper markdown text to explain your thought process and logic.

Proposed software solution shall have following components:

1. Instrument the data set – Exploratory Analysis
2. State or use prior work from academic / industry
3. Use appropriate state of the art AI algorithm (preferably deep learning based)
4. Enhance the data with synthetic images based on need and your algorithm (optional)
5. Reasonable accuracy on the given scope of time to develop
6. Rationales for the current Approach to enhance
7. Use any appropriate open source libraries and framework
8. Focus more on the algorithm than the overall application
9. Apply coding standards and best practices (We emphasize on clean code)
10. Make sure to give proper credits at the end (For example, citation for the dataset used in this challenge)

Delivery

- The solution shall be presented on GitHub / GitLab code with optional supported document in email.
- Execute all the cells of your notebook and export it to PDF. Include it in your source code repository.
- Prepare a 30 minute presentation on your solution describing challenges faced, choice of algorithms, results of your analysis and illustrations.

```
import numpy as np
import pandas as pd
import cv2
import matplotlib.pyplot as plt
from glob import glob
from tqdm.notebook import tqdm
import os
import shutil
import random
import json

from utils.general import get_coords, extract_from_txt, convert_to_yolov5, generate_json_coco, save_one_json,
from utils.plots import plot_one_box, plot_image_cropped_new

%matplotlib inline

Bad key "text.kerning_factor" on line 4 in
/home/jm/anaconda3/envs/yolo/lib/python3.7/site-packages/matplotlib/mpl-data/stylelib/_classic_test_patch.mplstyle.
You probably need to get an updated matplotlibrc file from
https://github.com/matplotlib/matplotlib/blob/v3.1.2/matplotlibrc.template
or from the matplotlib source distribution
```

Dataset exploratory analysis

The dataset is synthetic, meaning made of images that were rendered based on real models and is divided into 10 examples of cars and 4 classes, being infant seat, child seat, person and everyday object. To compose the dataset I have downloaded only the RGB images, and chose the first 6 cars, summing up a total of 12000 images. From those, 10%, 1200 images were randomly selected and moved to validation. The test set is already separated within each zip file, and it was added to the test folder.

The dataset is supposed to be separated into train, val and test, in a folder named images, inside the dataset folder, and the labels, which is in the xyxy format is supposed to be in a folder named labels_xyxy, inside the dataset folder, also in the train, val, test fashion, matching the names of the files '.png' and '.txt'.

```
# Support paths and label dictionary
image_path = './dataset/images/'
label_path = './dataset/labels_xyxy/'

labels_dict = {0: 'infant seat',
               1: 'child seat',
               2: 'person',
               3: 'everyday object'}

# All dataset files, for visualization
all_imgs = glob('./dataset/images/train/*.*')
all_labels = glob('./dataset/labels_xyxy/train/*.*')

# Get all the labels to check for quantity
label_hist = list()
for label in all_labels:
    info = get_coords(label)
    try:
        for l in info[:,0].tolist():
            label_hist.append(l-1)
    except:
        continue

# Plot labels histogram
plt.hist(label_hist)
plt.title('Labels Histogram')
plt.show()
```

Labels Histogram

In this Histogram we can see that we have more examples of person, and a lot less of everyday objects, which shows an imbalanced dataset.

```
# Get a random image from train set and plot to check
# Get random image and name
img_num = np.random.choice(len(all_imgs))
img_name = all_imgs[img_num].split('/')[-1].split('.')[0]

# Open the image with OpenCV
pic = cv2.imread(image_path + 'train/' + img_name + '.png')

# Get image box coordinates
box_coords_true = np.array(get_coords(label_path + 'train/' + img_name + '.txt'))
box_coords_true[:,0] = box_coords_true[:,0]-1

# Plot image
returned_true = plot_image_cropped_new(pic, box_coords_true)
plt.figure(figsize=(15,15))
plt.title('Image Preview: {}'.format(img_name))
plt.imshow(returned_true)
plt.show()
```

Image Preview: i3_train_imageID_1755_GT_4_0_2

State of the Art in Object Detection

With the advances in the usage of Transformers in Computer Vision, and Facebook AI have released Detectron2, prior works are the You Only Look Once (YOLO) series.

For this project, as I stated before, I decided to use the YOLO algorithm as I am familiar with and has a great performance and can be trained in smaller GPUs or even CPUs with the smallest model. I also have experience with the FasterRCNN model with a ResNet50 or mobilenetV2 backbone, but as it takes longer to converge, Yolo seems the best option, as also the available code is well maintained and complete, giving you several metrics while training and also image augmentations.

Yolo algorithm - Training the model with the Yolov5 repository

The algorithm expects the .yaml file inside the data to have the following informations:

```
train: ./dataset/images/train/
val: ./dataset/images/val/
test: ./dataset/images/test/

# number of classes
nc: 4

# class names
names: ["infant seat", "child seat", "person", "everyday object"]
```

Also the labels, had to be in the xywh format, instead of the xyxy that is the current annotation format.

```
# Preparing the dataset for the yolo requirements - xyxy to xywh

# Convert and save the annotations
task = 'train/'
all_txt = glob('./dataset/labels_xyxy/' + task + '*.txt')
for txt in tqdm(all_txt, desc='Converting Train'):
    info_dict = extract_from_txt(txt, task=task)
    convert_to_yolov5(info_dict, task)

# Convert and save the annotations
task = 'val/'
all_txt = glob('./dataset/labels_xyxy/' + task + '*.txt')
for txt in tqdm(all_txt, desc='Converting Val'):
    info_dict = extract_from_txt(txt, task=task)
    convert_to_yolov5(info_dict, task)

# Convert and save the annotations
task = 'test/'
all_txt = glob('./dataset/labels_xyxy/' + task + '*.txt')
for txt in tqdm(all_txt, desc='Converting Test'):
    info_dict = extract_from_txt(txt, task=task)
```

```
convert_to_yolov5(info_dict, task)

# Generate Json for COCO Eval - to be able to use the same metrics as it was used by the SVIRO team

# Annotations JSON generation
all_dataset = glob('./dataset/images/*/*.*')

json_dict = generate_json_coco(all_dataset)
```

```
anno_json = 'annotations.json' # predictions json
print(f'\nSaving {anno_json}...')
with open(anno_json, 'w') as f:
    json.dump(json_dict, f)
```

yolo, for 1 epoch. (I have trained a model in the Google Colab, to be able to use the GPU and train for more Epochs.)

The Jupyter notebook is also in the repository named Yolov5_bosh_colab_train.ipynb

```
!python train.py --img 640 --cfg models/yolov5s.yaml --hyp data/hyps/hyp.scratch.yaml --batch 16 --epochs 6 --
```

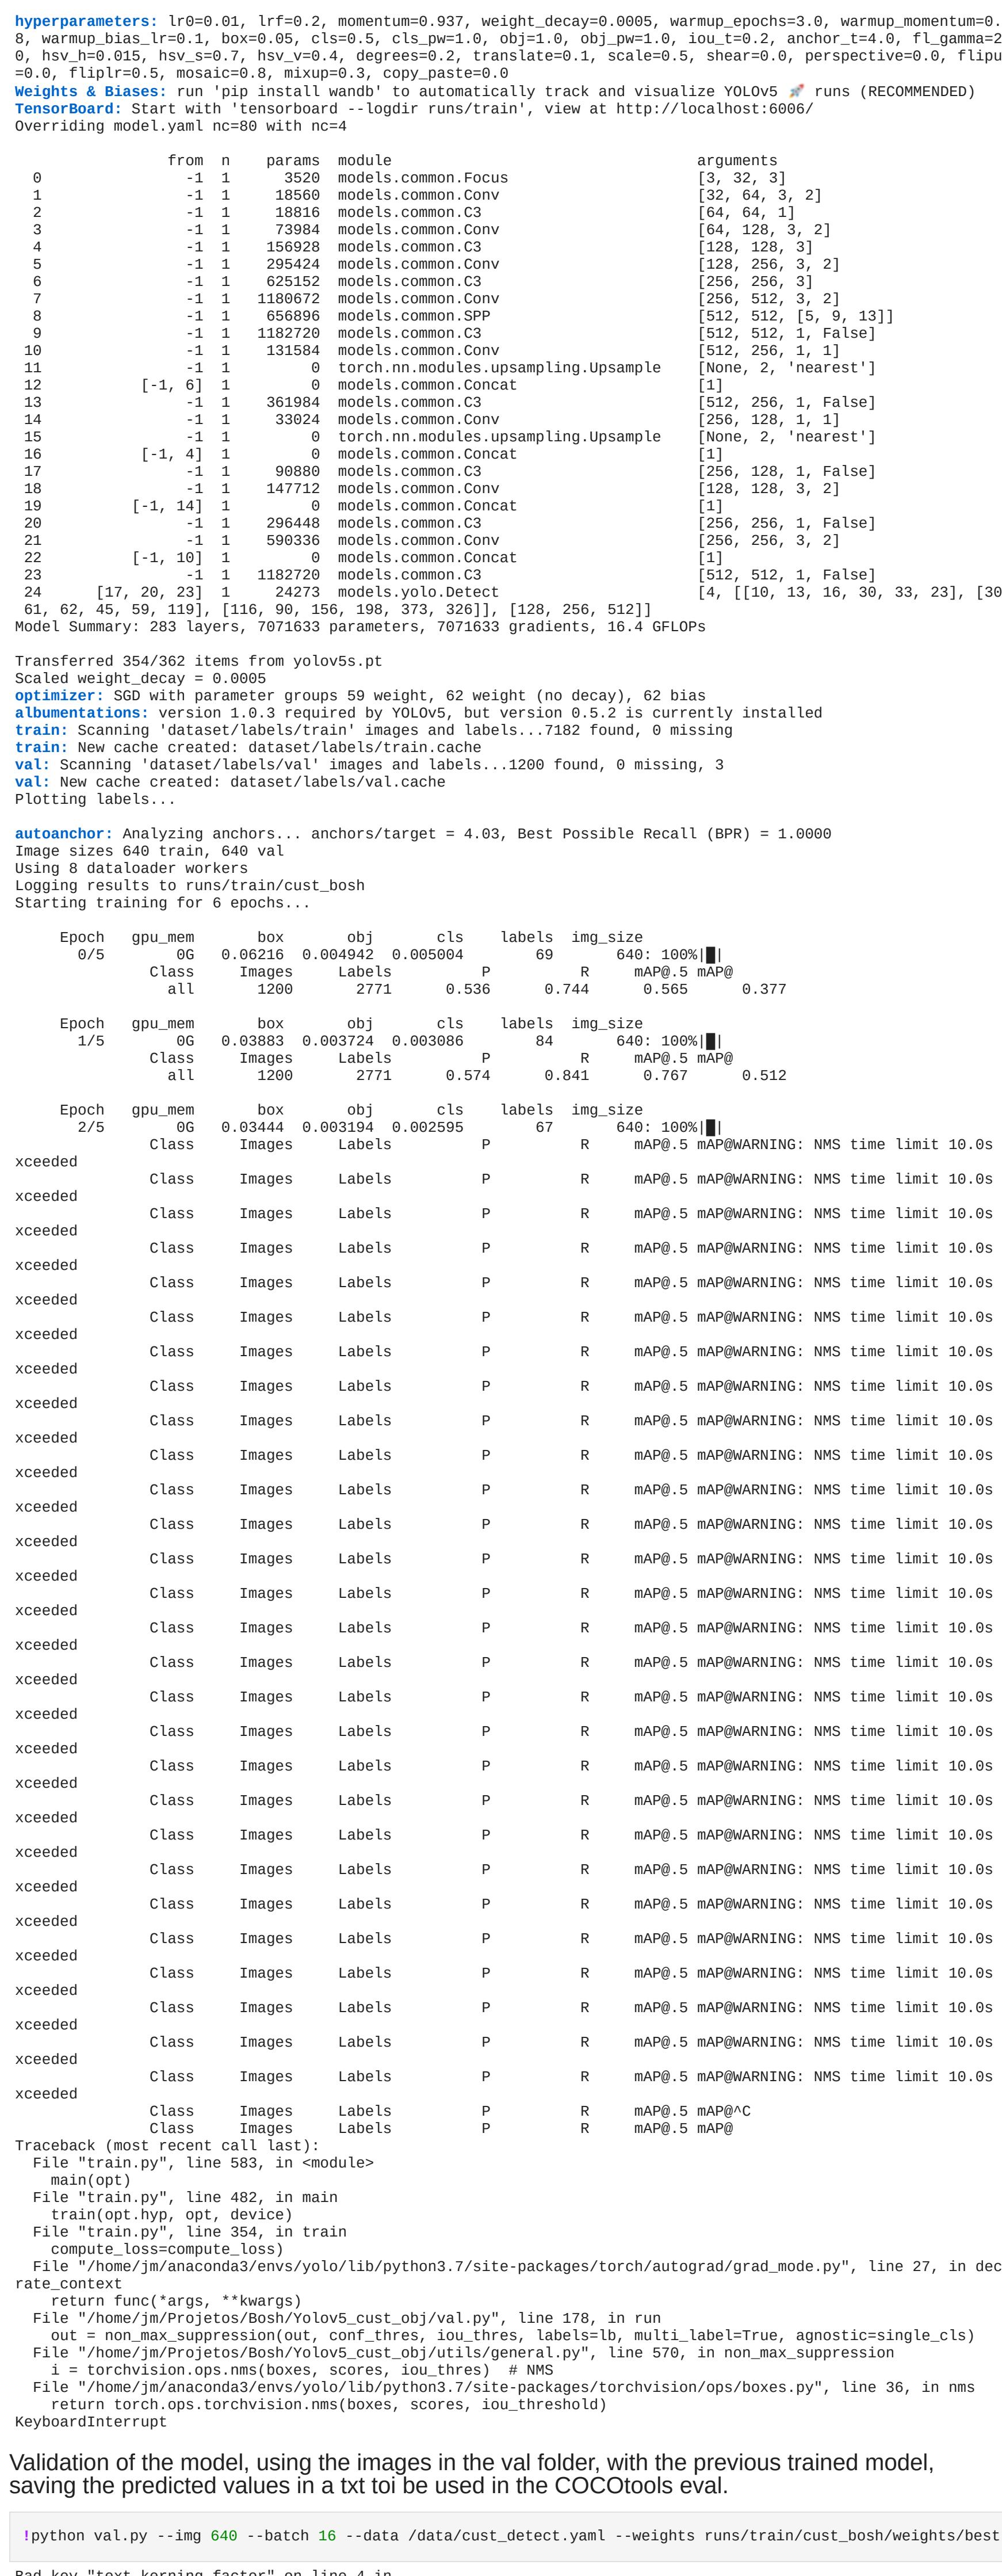
Bad key "text.kerning_factor" on line 4 in
/home/jm/anaconda3/envs/yolo/lib/python3.7/site-packages/matplotlib/mpl-data/stylelib/_classic_test_patch.mplstyle.

You probably need to get an updated `matplotlibrc` file from
<https://github.com/matplotlib/matplotlib/blob/v3.1.2/matplotlibrc.template>
or from the `matplotlib` source distribution

train: weights=yolov5s.pt, cfg=models/yolov5s.yaml, data=/data/cust_detect.yaml, hyp=data/hyps/hyp.scratch.yaml, epochs=6, batch_size=16, imgsz=640, rect=False, resume=False, nosave=False, noval=False, noautoanchor=False, evolve=None, bucket=, cache_images=False, image_weights=False, device=, multi_scale=False, single_cls=False, dam=False, sync_bn=False, workers=8, project=runs/train, entity=None, name=cust_bosh, exist_ok=False, quad=False, linear_lr=False, label_smoothing=0.0, upload_dataset=False, bbox_interval=-1, save_period=-1, artifact_alias=latest, local_rank=-1

github: skipping check (not a git repository), for updates see <https://github.com/ultralytics/yolov5>

YOLOv5 🚀 2021-7-28 torch 1.8.1+cu102 CPU



```
Bad key   text.kerning_factor' on b
/home/jm/anaconda3/envs/yolo/lib/p
yle.
You probably need to get an update
https://github.com/matplotlib/matpl
or from the matplotlib source dist
val: data=../data/cust_detect.yaml,
       conf_thres=0.001, iou_thres=0.6,
       ruse, save_hybrid=False, save_conf=
False
```

```
Fusing layers...
Model Summary: 224 layers, 7062001 parameters, 0 gradients, 16.4 GFLOPS
val: Scanning 'dataset/labels/val.cache' images and labels... 1200 found, 0 miss
      Class    Images     Labels       P       R   mAP@.5   mAP@
        all     1200     2771     0.574     0.841     0.767     0.513
  infant seat     1200      553     0.718     0.774     0.784     0.49
 child seat     1200      704     0.401     0.986     0.759     0.591
      person     1200     1110     0.56      0.965     0.914     0.62
everyday object     1200      404     0.619     0.639     0.61      0.351
Speed: 0.5ms pre-process, 90.3ms inference, 176.0ms NMS per image at shape (16, 3, 640, 640)
Results saved to runs/val/val_bosh2
1200 labels saved to runs/val/val_bosh2/labels
```

```
Yolov5 🚀 2021-7-28 torch 1.8.1+cu102 CPU

Fusing layers...
Model Summary: 224 layers, 7062001 parameters, 0 gr
test: Scanning 'dataset/labels/test' images and lab
test: New cache created: dataset/labels/test.cache
```

```
    all        1500      3536      0.723      0.785      0.846      0.655
  infant seat     1500       725      0.85      0.578      0.815      0.607
   child seat     1500       841      0.734      0.751      0.779      0.612
      person      1500      1409      0.832      0.933      0.958      0.74
everyday object    1500       561      0.476      0.879      0.832      0.659
Speed: 0.6ms pre-process, 86.2ms inference, 308.7ms NMS per image at shape (16, 3, 640, 640)
Results saved to runs/val/test_bosh_colab
1500 labels saved to runs/val/test_bosh_colab/labels
```

```
# returns JSON object as
# a dictionary
data_json = json.load(f)

# Closing file
f.close()

jdict = []
all_ids = []
```

```
for files in txt_files:
    l_list = list()
    box_coords_pred = list()
    r_labels = pd.read_table(files, names=['infos']).values.tolist()
    for b in r_labels:
        label, x, y, w, h, score = list(map(float, b[0].split()))
        x1, y1, x2, y2 = unconvert(640, 640, x, y, w, h)
        l_list.append([label, x1, y1, x2, y2, score])
        if score > 0.5:
            box_coords_pred.append([label, x1, y1, x2, y2])
```

```
save_one_json(np.array(l_list), jdict, files, data_json, all_ids)

from pycocotools.coco import COCO
from pycocotools.cocoeval import COCOeval

if True and len(jdict):
    pred_json = 'predictions.json' # predictions json
    print(f'\nEvaluating pycocotools mAP... saving {pred_json}...')
    with open(pred_json, 'w') as f:
```

```
    json.dump(jdict, f)

anno_json = 'annotations.json' # predictions json
print(f'\nEvaluating pycocotools mAP... saving {anno_json}...')
try: # https://github.com/cocodataset/cocoapi/blob/master/PythonAPI/pycocoEvalDemo.ipynb
    anno = COCO(anno_json) # init annotations api
    pred = anno.loadRes(pred_json) # init predictions api
    eval = COCOeval(anno, pred, 'bbox')
    eval.params.imgIds = all_ids # image IDs to evaluate
    eval.evaluate()
    eval.accumulate()
    eval.summarize()
```

```
    map_, map50 = eval.stats[:2] # update results (mAP@0.5:0.95, mAP@0.5)
except Exception as e:
    print(f'pycocotools unable to run: {e}')
```

Evaluating pycocotools mAP... saving predictions.json...

Evaluating pycocotools mAP... saving annotations.json...

loading annotations into memory...

Done (t=0.06s)

creating index...

index created!

```
index created!
Loading and preparing results...
DONE (t=2.19s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=7.34s).
Accumulating evaluation results...
DONE (t=1.24s).
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.217
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.286
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.235
```

Average Precision	(AP) @ [IoU=0.75]	area= all	maxDets=100	= 0.233
Average Precision	(AP) @ [IoU=0.50:0.95]	area= small	maxDets=100	= -1.000
Average Precision	(AP) @ [IoU=0.50:0.95]	area=medium	maxDets=100	= -1.000
Average Precision	(AP) @ [IoU=0.50:0.95]	area= large	maxDets=100	= 0.218
Average Recall	(AR) @ [IoU=0.50:0.95]	area= all	maxDets= 1	= 0.396
Average Recall	(AR) @ [IoU=0.50:0.95]	area= all	maxDets= 10	= 0.870
Average Recall	(AR) @ [IoU=0.50:0.95]	area= all	maxDets=100	= 0.878
Average Recall	(AR) @ [IoU=0.50:0.95]	area= small	maxDets=100	= -1.000
Average Recall	(AR) @ [IoU=0.50:0.95]	area=medium	maxDets=100	= -1.000
Average Recall	(AR) @ [IoU=0.50:0.95]	area= large	maxDets=100	= 0.878

```
# Visualizing some predicted bounding boxes with the respective image

# Visualizing the results

# All labels files, for visualization
all_labels = glob('./runs/val/test_bosh_colab/labels/*.*')

# Get a random image from train set and plot to check

# Get random image and name
image_name = random.choice(os.listdir('train'))
```

```
img_num = np.random.choice(len(all_labels))
img_name = all_labels[img_num].split('/')[-1].split('.')[0]

# Open the image with OpenCV
pic = cv2.imread(image_path + 'test/' + img_name + '.png')
pic = cv2.resize(pic, (640, 640))

# Get image box coordinates

label_pred = pd.read_table('./runs/val/test_bosh_colab/labels/' + img_name + '.txt', names=['cls'])
box_coords_pred = list()
```

```
for box in label_pred['cls']:
    label, x, y, w, h, score = list(map(float, box.split()))
    x1, y1, x2, y2 = unconvert(640, 640, x, y, w, h)
    if score > 0.3:
        box_coords_pred.append([label, x1, y1, x2, y2])

# Plot image
returned_true = plot_image_cropped_new(pic, box_coords_pred)
plt.figure(figsize=(15,15))
plt.title('Image Preview: {}'.format(img_name))
```

```
plt.imshow(returned_true)  
plt.show()
```

Image Preview: aclass_test_imageID_167_GT_6_0_4

A close-up view of a dark, textured surface, possibly the interior of a car. A bright green rectangular box is overlaid on the image, containing the white text "child seat". The text is centered within the green box. The background is dark and grainy, suggesting a low-light environment or a close-up shot of a textured material.

A dark, abstract image featuring a bright green vertical stripe on the left and a white text overlay "everyday objec..." on the right.

This image shows a dark, textured surface, likely a wall or floor. A thick vertical green line runs along the left edge. In the upper right quadrant, there is a small, rectangular blue object with a black frame around it.

Confusion Matrix generated by the yolo output